



CSU33031 Computer Networks

File Retrieval Protocol

September 11, 2022

1 Introduction

The focus of this assignment is to learn about protocol development and the information that is kept in a header to support functionality of a protocol. Protocol design is always a balancing act between introducing functionality that relies on additional header information and the overhead that the additional header information introduces.

2 Protocol Details

The aim of the protocol is to provide a mechanism to retrieve files from a service based on UDP datagrams. The encoding of the header information of this protocol should be implemented in a binary format. The protocol involves a number of actors: One or more clients, an ingress node, and one or more workers. A client issues requests for files to an ingress node and receives replies from this node. The ingress node processes requests, forwards them to one of the workers that are associated with it, and forwards replies to clients that have sent them. The header information that you included in your packets has to support the identification of the requested action, the transfer of files - potentially consisting of a number of packets and the management of the workers by the server.

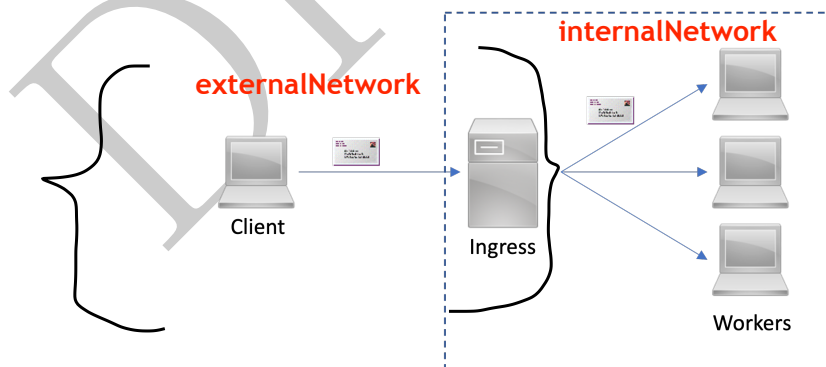


Figure 1: A client sends a request to a server, called ingress here, which distributes requests to available workers.

The basic functionality that the file request protocol has to provide is to support requests for files from a server which then distributes the requests to workers. The description of the scenario above shows that there are a number of additional functionalities that a protocol may need to provide such as various file sizes, multiple concurrent streams, security, etc. The QUIC Protocol [7, 5] is an example of a protocol that you could look at for functionalities that may be interesting for you to include in your protocol. In the end, it is

your decision which functionalities you implement in your protocol. In the deliverables, you need to describe these functionalities and justify why you included them in your protocol.

The following flow diagrams, figure 2 is an example of visualizations of network traffic between components. They show the sequence of messages exchanged between components with the start of their transmission at the sender and their arrival at the receiver.

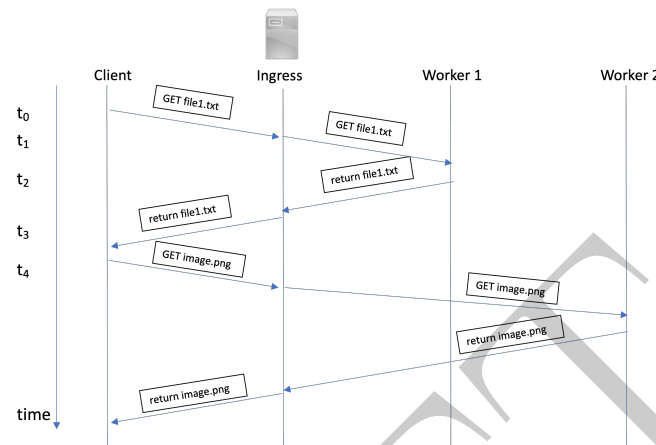


Figure 2: A client will request a file from the server. The server will select a worker and forward the request to the worker. The worker will retrieve the file and send it to the server which then will forward it to the client. The client will then request another file from the server and the server will pick a different worker to return this specific file.

Please use Flow Diagrams like these to document the communication between components of your implementation in your videos and in your final report.

The protocol can be implemented in a programming language of your choosing. One of the conditions is that the communication between the processes is realized using **UDP sockets and UDP Datagrams**. Please avoid Python 2.7 because the implementation of Datagram sockets in the obsolete versions is based the transfer of strings instead of binary arrays.

The easiest way to start with the development of your solution is possibly to connect your components through the localhost interface of a machine; however, at the end, you will need to be able to demonstrate that your protocol can connect components located at a number of hosts. There are a number of platforms that support the simulation of topologies or provide virtual infrastructures e.g. Docker [2], Mininet [6], Kubernetes [1], etc. For someone starting with socket programming and networking, I would suggest to use a platform such as Docker or Mininet; for someone already familiar with these concepts, I would suggest to implement their solution using Kubernetes. However, these are only suggestion and you need to make the decision how to implement your solution.

3 Deliverables & Submission Details

The deliverables for this assignment are split into 3 parts: 2 videos, a report describing your solution and the files making up the implementation of your solution. The deadline for the submission of these deliverables are given in Blackboard.

One component of the deliverables at every step is the submission of captures of network traffic. These captures should be in the form of PCAP files [3]. Programs such as Wireshark [4], tshark, etc offer functionality to capture network traffic from interfaces

3.1 Part 1: Video & PCAP file

The video of part 1 should demonstrate the initial design of your solution and a capture of network traffic between the components of your solution and the files of your traffic capture in pcap format. In the video,

you should explain the setup of the topology that you are using and the information that makes up the header information in your traffic captures.

The submission process for this part consists of two steps: 1) Submitting the PCAP file or files that you captured from your network traffic, and 2) submitting the video for this step.

The video is to be no longer than 4 minutes; content past the 4 minute-mark will be ignored during marking. Videos with voice-over using text-to-speech (TTS) will not be accepted and marked with 0.

3.2 Part 2: Video & PCAP file

The video of part 2 should demonstrate the current state of your solution, the functionality that you have implemented so far, and a capture of network traffic between the components of your solution and the files of your traffic capture in pcap format. In the video, you should explain the basic implementation of your protocol and the information that is being exchanged between the components of your solution.

The submission process for this part consists of two steps: 1) Submitting the PCAP file or files that you captured from your network traffic, and 2) submitting the video for this step.

Videos with voice-over using text-to-speech (TTS) will not be accepted and marked with 0. The video is to be no longer than 4 minutes; content past the 4 minute-mark may be ignored during marking.

3.3 Final Deliverable

The final deliverable should include a report that describes the components of your solution and their functionality, the protocol that you implemented and the communication between the components of your solution, the topology that you used to run your solution and how your solution was executed.

The submission process for this part consists of three steps: 1) Submitting the PCAP file or files that you captured from your network traffic, 2) submitting the source code and any files that may be necessary to execute your solution, and 3) submitting the report about your solution.

The files that contain the implementation and the report should be submitted through Blackboard. Every file should contain the name of the author. The source files of the implementation should be submitted as an archived file e.g. ".zip" or ".tar.gz". The report should be submitted as either word- or pdf-document. The deadline for the submission is given in Blackboard.

The report may be submitted to services such as TurnItIn for plagiarism checks.

4 Marking Scheme

The contribution of the assignment to the overall mark for the module is 30% or 30 points. The submission for part 1 and 2 will be each marked out of 5 points and the submission for the final part will be marked out of 20 points. The mark for the final deliverable will be split into 50% for the functionality of your solution and 50% for the documentation through the report.

References

- [1] The Kubernetes Authors. Kubernetes project page. <https://kubernetes.io>, visited Sep 2021.
- [2] Docker. Docker project page. <https://www.docker.com>, visited Sep 2021.
- [3] Wikipedia Editors. pcap - wikipedia page. <https://en.wikipedia.org/wiki/Pcap>, visited Aug 2021.
- [4] Wireshark Foundation. Wireshark project page. <https://www.wireshark.org>, visited Sep 2021.
- [5] Adam Langley, Alistair Riddoch, Alyssa Wilk, Antonio Vicente, Charles Krasnic, Dan Zhang, Fan Yang, Fedor Kouranov, Ian Swett, Janardhan Iyengar, Jeff Bailey, Jeremy Dorfman, Jim Roskind, Joanna Kulik, Patrik Westin, Raman Tenneti, Robbie Shade, Ryan Hamilton, Victor Vasiliev, Wan-Teh Chang, and Zhongyi Shi. The quic transport protocol: Design and internet-scale deployment. In *Proceedings of the Conference of the ACM Special Interest Group on Data Communication, SIGCOMM '17*, page 183–196, New York, NY, USA, 2017. Association for Computing Machinery.

- [6] Mininet. Mininet project page. <http://mininet.org>, visited Sep 2021.
- [7] The Chromium Projects. Quic, a multiplexed transport over udp. <https://www.chromium.org/quic/>, visited Sep 2022.

DRAFT